

Model #6

Ethan Allavarpu (UID: 405287603)

10/28/2020

Transforming and Cleaning the Data

```
training <- read.csv("training.csv", stringsAsFactors = TRUE)
training$class <- factor(training$class)
levels(training$class) <- c("NG", "OG", "TSG")
outlier <- function(data) {
  low <- mean(data) - 3 * sd(data)
  high <- mean(data) + 3 * sd(data)
  which(data < low | data > high)
}
library(ggplot2)
scatter <- function(var) {
  ggplot(training, aes_string(var, "class")) +
    geom_jitter(width = 0.05, height = 0.1, size = 0.1,
               colour = rgb(0, 0, 0, alpha = 1 / 3))
}
scat_plot <- lapply(names(training)[-99], scatter)
library(gridExtra)
# grid.arrange(grobs = scat_plot[1:20], ncol = 4)
# grid.arrange(grobs = scat_plot[21:40], ncol = 4)
# grid.arrange(grobs = scat_plot[41:60], ncol = 4)
# grid.arrange(grobs = scat_plot[61:80], ncol = 4)
# grid.arrange(grobs = scat_plot[81:98], ncol = 4)
outlier_index <- sort(table(unlist(lapply(training[, -99], outlier))), decreasing = TRUE)
outlier_index[1:100]
```

915	1280	2918	517	1914	2182	3052	1173	2215	3049	259	740	1749	1979	2998	417
24	24	24	22	22	22	20	19	19	19	18	18	18	18	18	17
441	806	2297	422	635	1258	1570	2278	2518	2729	80	150	2694	169	276	341
17	17	17	16	16	16	16	16	16	16	15	15	15	14	14	14
1528	1556	1726	1809	1911	1955	2071	2624	2641	3120	3142	73	277	364	751	1244
14	14	14	14	14	14	14	14	14	14	14	13	13	13	13	13
1330	2329	2787	343	1138	1171	1188	1372	1460	2031	2251	2968	2983	3166	352	634
13	13	13	12	12	12	12	12	12	12	12	12	12	12	11	11
907	923	1096	1858	2636	588	1137	1317	1463	1561	1740	1991	2487	2540	2555	2621
11	11	11	11	11	10	10	10	10	10	10	10	10	10	10	10
2815	3029	74	144	657	789	857	1267	1610	1932	2022	2093	2142	2534	2666	2721
10	10	9	9	9	9	9	9	9	9	9	9	9	9	9	9
2848	2900	3027	155												
9	9	9	8												

```

training <- training[-as.numeric(names(outlier_index)[1:50]),]
sort(training$Missense_TO_Silent_Ratio, decreasing = TRUE)[1:10]

[1] 384.98658 172.91420 135.59623 71.09712 23.21809 21.81193 20.37791
[8] 19.42402 19.38769 15.84808

training <- training[-which(training$Missense_TO_Silent_Ratio > 100), ]
sort(training$Missense_KB_Ratio, decreasing = TRUE)[1:10]

[1] 2063.9413 1296.6625 1060.0601 952.3810 931.4227 726.8519 594.7603
[8] 593.3610 581.5085 516.8084

training <- training[-which(training$Missense_KB_Ratio > 2000), ]
sort(training$Lof_TO_Silent_Ratio, decreasing = TRUE)[1:10]

[1] 81.177835 9.030120 6.470238 5.582840 4.741460 4.558252 4.176630
[8] 4.058140 4.039062 4.021930

training <- training[-which(training$Lof_TO_Silent_Ratio > 5), ]
sort(training$Gene_expression_Z_score, decreasing = TRUE)[1:10]

[1] 19.720 9.210 7.080 6.883 6.590 6.280 5.321 5.316 3.161 2.767

training <- training[-which(training$Gene_expression_Z_score > 4), ]
sort(training$dN_to_dS_ratio, decreasing = TRUE)[1:10]

[1] 20.950 3.649 3.446 3.372 2.574 2.194 2.183 2.102 1.921 1.744

training <- training[-which(training$dN_to_dS_ratio > 5),]
sort(training$Silent_KB_Ratio, decreasing = TRUE)[1:10]

[1] 474.4745 193.1684 174.0558 171.0362 166.4971 160.2273 158.7697 148.5800
[9] 143.6782 135.2657

training <- training[-which(training$Silent_KB_Ratio > 200), ]
sort(training$Lost_start_and_stop_fraction, decreasing = TRUE)[1:10]

[1] 0.333 0.167 0.118 0.087 0.074 0.071 0.071 0.068 0.067 0.067

training <- training[-which(training$Lost_start_and_stop_fraction > 0.2),]
sort(training$Synonymous_Zscore, decreasing = FALSE)[1:10]

[1] -20.5110 -10.9780 -10.2960 -9.7346 -9.3720 -8.8090 -8.4062 -8.3918
[9] -8.1076 -8.1076

training <- training[-which(training$Synonymous_Zscore < -15), ]
numeric_training <- training[,-99]

n_zeroes <- rep(NA, nrow(numeric_training))

for(i in seq_len(nrow(numeric_training))){
  row_i_zeroes <- 0
  for(j in seq_len(ncol(numeric_training))){
    if(round(numeric_training[i,j], digits = 5) == 0){
      row_i_zeroes <- row_i_zeroes + 1
    }
  }
  n_zeroes[i] <- row_i_zeroes
}

```

```
training <- training[n_zeroes <= 50, ]
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following object is masked from 'package:gridExtra':

combine

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
#function to calculate wca
```

```
score <- function (conf_mat) {  
  print(sum(diag(conf_mat) * c(1, 20, 20)))  
  print(sum(diag(conf_mat) * c(1, 20, 20)) / sum(apply(conf_mat, 2, sum) * c(1, 20, 20)))  
}
```

```
# Set new threshold to account for unbalanced data
```

```
classify <- function(probs) {  
  if (any(probs[2:3] > 0.05)) {  
    subset <- probs[2:3]  
    output <- which(subset == max(subset))  
    if (length(output) > 1) {  
      output <- sample(1:2, 1)  
    }  
  } else {  
    output <- 0  
  }  
  output  
}
```

Multinom (Logistic Regression)

```
library(dplyr)
library(caret)
```

Loading required package: lattice

```
set.seed(43)
vars <- training %>% select(Broad_H3K9ac_percentage, N_LOF, pLOF_Zscore,
                           Missense_Entropy,
                           N_Splice, LOF_TO_Total_Ratio, VEST_score,
                           BioGRID_log_degree,
                           Broad_H3K79me2_percentage, FamilyMemberCount,
                           S50_score_replication_timing, Gene_expression_Z_score,
                           Polyphen2, Broad_H3K36me3_percentage, class)
vars_test <- createDataPartition(vars$class, p = 0.76,
                                  list = FALSE)
vars_train <- vars[vars_test, ]
vars_test <- vars[-vars_test, ]

mn <- nnet::multinom(class ~ ., data = vars_train, model = TRUE)

# weights: 48 (30 variable)
initial value 2529.005489
iter 10 value 1586.420047
iter 20 value 903.363165
iter 30 value 480.938933
iter 40 value 365.932995
iter 50 value 363.703288
iter 60 value 363.693907
final value 363.693899
converged

tidymn <- broom::tidy(mn) %>% arrange(p.value)
terms <- tidymn$term[-(1:2)]
terms <- unique(terms)

vars_mn <- training %>% select("VEST_score", "BioGRID_log_degree",
                              "Missense_Entropy",
                              "Broad_H3K9ac_percentage",
                              "N_LOF", "Broad_H3K36me3_percentage",
                              "pLOF_Zscore", "N_Splice", "LOF_TO_Total_Ratio",
                              "FamilyMemberCount", "Polyphen2",
                              "Broad_H3K79me2_percentage", "Gene_expression_Z_score",
                              "S50_score_replication_timing"
                              ) # these are just the predictors it liked

tests <- read.csv("test.csv")
preds <- predict(mn, newdata = vars_test, type = "prob")
predclass <- apply(preds, 1, classify)
tbl <- table(predclass, vars_test$class)
score(tbl)
```

```
[1] 1538
[1] 0.7767677
```

tbl

predclass	NG	OG	TSG
0	598	3	3
1	33	24	5
2	29	8	23

```
test_preds <- predict(mn, newdata = tests, type = "prob")
test_preds <- apply(test_preds, 1, classify)
csv_file <- data.frame("id" = tests$id,
                      "class" = test_preds)
write.csv(csv_file, "modelpredictions7_LOG.csv", row.names = FALSE)
```